



Chapter 8

Exercise 1:

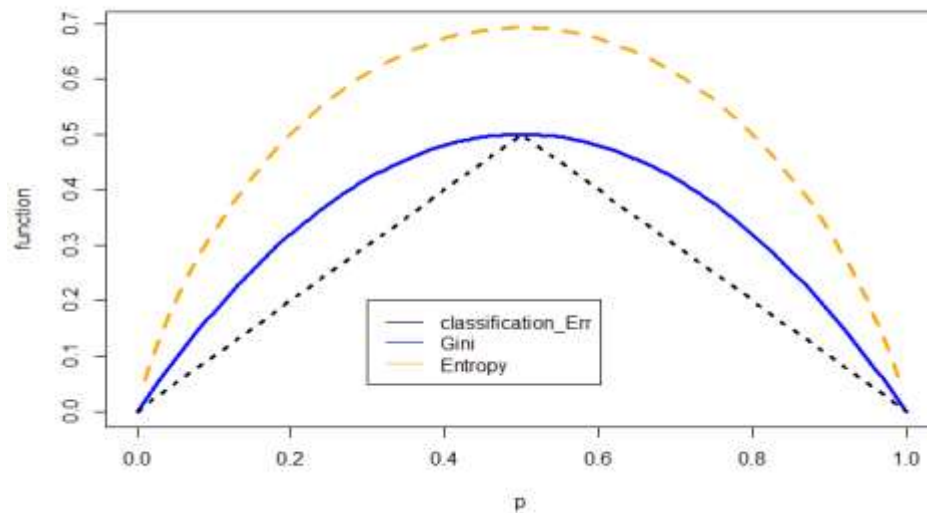


$$\begin{array}{lll}
 t_1 = 100 & t_3 = 150 & t_5 = 180 \\
 t_2 = 40 & t_4 = 50 &
 \end{array}$$



Exercise 3:

```
> p=seq(0,1,0.01)
> gini= p * (1-p) * 2
> entropy= -(p * log(p) + (1-p) * log(1-p))
> classification_error= 1 - pmax(p,1-p)
> matplot(p,cbind(gini,entropy,classification_error),type = "l",col = c("blue","orange","black"),xlab = "p",ylab = "function of
'p'",lwd = 3)
>
```




Exercise 4:

a) Sketch the tree corresponding to the partition of the predictor space illustrated in the left-hand panel of Figure 8.14. The numbers inside the boxes indicate the mean of Y within each region


Chapter 8

Exercise 4

a)



b) Create a diagram similar to the left-hand panel of Figure 8.14, using the tree illustrated in the right-hand panel of the same figure. You should divide up the predictor space into the correct regions and indicate the mean for each region.



Exercise 5:

Suppose we produce ten bootstrapped samples from a data set containing red and green classes. We then apply a classification tree to each bootstrapped sample and, for a specific value of X , produce 10 estimates of $P(\text{Class is Red} | X)$: 0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, and 0.75. There are two common ways to combine these results together into a single class prediction. One is the majority vote approach discussed in this chapter. The second approach is to classify based on the average probability. In this example, what is the final classification under each of these two approaches?

Solution:

The majority vote approach classifies X as red, from given 10 estimates 4 estimates have a probability less than 0.5 and 6 estimates have a probability > 0.5 .

The average Probability Approach classifies X as green as an average of 10 estimates is less than 0.5

$$\begin{aligned} \text{Avg} &= (0.1 + 0.15 + 0.2 + 0.2 + 0.55 + 0.6 + 0.6 + 0.65 + 0.7 + 0.75) / 10 \\ &= 0.45 \end{aligned}$$


Since the Avg < 0.5 , the Class for X will be Green.

Chapter 9

Exercise 1:


Sketch the hyperplane $1 + 3X_1 - X_2 = 0$. Indicate the set of points for which $1 + 3X_1 - X_2 > 0$, as well as the set of points for which $1 + 3X_1 - X_2 < 0$.

```
> x1= seq(-1,1,0.1)
> x2= 1+3 *x1
> plot(x1,x2,xlab="x1",ylab="x2",type="l",lwd=3,col="blue")
> text(0,3,expression(1+3*x1-x2<0),col="black")
> text(-0,-1,expression(1+3*x1-x2 > 0),col="black")
> for(i in seq(-1,1,length.out = 5)){
+   pts=data.frame(rep(i,5),seq(-2,4,length.out = 5))
+   points(pts,col=ifelse(1+3*pts[,1]-pts[,2]>0,'red','green'),bg="grey")
+ }
> |
```



B. On the same plot, sketch the hyperplane $-2 + X_1 + 2X_2 = 0$. Indicate the set of points for which $-2 + X_1 + 2X_2 > 0$, as well as the set of points for which $-2 + X_1 + 2X_2 < 0$

```
+ }
> lines(x1,1-x1/2,col="purple",lwd=3)
> text(-0.5,2,expression(-2 +x1+2*x2<0),col="black")
> text(0.5,0,expression(-2 +x1+2*x2>0),col="black")
> for(i in seq(-1,1,length.out = 5)){
+   pts=data.frame(rep(i,5),seq(-2,4,length.out = 5))
+   points(pts,col=ifelse(1+3*pts[,1]-pts[,2]>0,'red','green'),pch=ifelse(-2+pts[,1]+2*pts[,2]>0,1,2))
+ }
> |
```




Exercise 2:

We have seen that in $p = 2$ dimensions, a linear decision boundary takes the form $\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$. We now investigate a non-linear decision boundary.

- a) Sketch the curve $(1 + X_1)^2 + (2 - X_2)^2 = 4$.


The equation describes a circle centered at the point $(-1, 2)$

```
> plot(NA, NA, type = "n", xlim = c(-4, 2), ylim = c(-1, 5), asp = 1, xlab = "x1", ylab = "x2")
> symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE, lwd=3)
```




- b) On your sketch, indicate the set of points for which $(1 + X_1)^2 + (2 - X_2)^2 > 4$, as well as the set of points for which $(1 + X_1)^2 + (2 - X_2)^2 \leq 4$

```
> plot(NA, NA, type = "n", xlim = c(-4, 2), ylim = c(-1, 5), asp = 1, xlab = "x1", ylab = "x2")
> symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
> text(c(-1), c(2), expression((1+x1)^2 + (2-x2)^2 <= 4), col = "red")
> text(c(-6), c(2), expression((1+x1)^2 + (2-x2)^2 > 4), col = "red")
>
```



- c) Suppose that a classifier assigns an observation to the blue class if $(1 + X_1)^2 + (2 - X_2)^2 > 4$, and to the red class otherwise. To what class is the observation (0, 0) classified? (-1, 1)? (2, 2)? (3, 8)?

```
> plot(NA, NA, type = "n", xlim = c(-4, 2), ylim = c(-1, 5), asp = 1, xlab = "x1", ylab = "x2")
> symbols(c(-1), c(2), circles = c(2), add = TRUE, inches = FALSE)
> text(c(-1), c(2), expression((1+x1)^2 + (2-x2)^2 <= 4), col = "red")
> text(c(-6), c(2), expression((1+x1)^2 + (2-x2)^2 > 4), col = "red")
> for(i in seq(-6, 4, length.out = 10)){
+   pts=data.frame(rep(i, 10), seq(-1, 5, length.out = 10))
+   points(pts, col=ifelse( (1 + pts[,1])^2 + (2-pts[,2])^2 > 4, 'blue', 'red'))
+ }
> |
```



We can see from the graph that points (0,0), (2,2), and (3,8) belong to the class blue so classify them as blue, and (-1,1) classify as red.

We can also predict class by the substitute value of coordinate (x, y) in a given equation and to check if the result is less or greater than 4.

- d) Argue that while the decision boundary in (c) is not linear in terms of X_1 and X_2 , it is linear in terms of X_1 , X_1^2 , X_2 , and X_2^2


The decision boundary is simply the boundary described by $f(X)=0$

Means $(1 + X_1)^2 + (2 - X_2)^2 - 4 = 0 \rightarrow x_1^2 + x_2^2 + 2x_1 - 4x_2 + 1 = 0$ which is linear in term of x_1, x_1^2, x_2 and x_2^2 .

Exercise 3:

We are given $n = 7$ observations in $p = 2$ dimensions. For each observation, there is an associated class label.

```
> x1 = c(3, 2, 4, 1, 2, 4, 4)
> x2 = c(4, 2, 4, 4, 1, 3, 1)
> colors = c("red", "red", "red", "red", "blue", "blue", "blue")
> plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5), lwd=3)
> |
```




- b) Sketch the optimal separating hyperplane, and provide the equation for this hyperplane (of the form (9.1))

We can see in the plot, the optimal separating hyperplane must be between observations (2,1) and (2,2), and between observations (4,3) and (4,4). So, it is a line that passes through the points (2,1.5) and (4,3.5) whose equation is

$$x_1 - x_2 - 0.5 = 0$$

```
plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5), lwd=3)
abline(-0.5, 1, col="green", lwd=3)
|
```



c) Describe the classification rule for the maximal margin classifier. It should be something along the lines of "Classify to Red if $\beta_0 + \beta_1 x_1 + \beta_2 x_2 > 0$, and classify to Blue otherwise." Provide the values for β_0 , β_1 , and β_2 .

Line Equation is $x_1 - x_2 - 0.5 = 0 \rightarrow x_2 = x_1 - 0.5$

We can observe that the maximal margin classifier will be red if $x_2 > x_1 - 0.5$ else the classifier will be blue.

Maximum margin classifier $f(x) = x_2 - x_1 + 0.5$

$f(x) > 0 \rightarrow$ observation class is red


$f(x) \leq 0 \rightarrow$ observation class is Blue.

We therefore have coefficients = (0.5, -1, 1)

d) On your sketch, indicate the margin for the maximal margin hyperplane.

The solid line indicates the maximal margin hyperplane, and the margin is the distance from the solid green line to either of the blue lines

```
> plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5),lwd=3)
> abline(-0.5,1,col="green",lwd=3)
> abline(-1,1,col="blue",lwd=2)
> abline(0,1,col="blue",lwd=2)
> |
```



e) Indicate the support vectors for the maximal margin classifier

The support vectors for the maximal margin classifier are $(2,1)$, $(2,2)$, $(4,3)$, and $(4,4)$


f) Argue that a slight movement of the seventh observation would not affect the maximal margin hyperplane.

From the above graph, we can see that the observation $(4,1)$ is not a support vector. So, any movement of his observation in any direction would not affect the maximum margin hyperplane.

The observation $(4,1)$ would have to be inside the margin to start impacting the position of the maximal margin hyperplane.

g) Sketch a hyperplane that is not the optimal separating hyperplane, and provide the equation for this hyperplane

The hyperplane given by the line is $x_1 - x_2 - 0.25 = 0$




As we can see from the graph the line equation above separates all observations, but it is not optimal as its margin is smaller than the optimal margin.

(h) Draw an additional observation on the plot so that the two classes are no longer separable by a hyperplane.

If we add an observation of either red point (2,3) or blue point (3,1) to the plot, make the two classes no longer be separated by a hyperplane.

```
> plot(x1, x2, col = colors, xlim = c(0, 5), ylim = c(0, 5), lwd=3)
> points(c(3),c(1),col=c("red"),lwd=3)
> points(c(2),c(3),col=c("blue"),lwd=3)
>
```



Problem 2.1

Simulate a binary classification dataset with a single feature via a mixture of normal distributions using R. The normal distribution parameters (using the function `rnorm`) should be (5,2) and (-5,2) for the pair of samples.

```
> library(rpart)
> library(rpart.plot)
> set.seed(200)
> d1=data.frame(val=rnorm(n=200,mean=5,sd=2),y=rep("yes",200))
> d2=data.frame(val=rnorm(n=200,mean=-5,sd=2),y=rep("no",200))
> data1=rbind(d1,d2)
> summary(data1)
```

	val	y
Min.	:-10.43362	Length:400
1st Qu.	:-5.06981	Class :character
Median	: 1.20257	Mode :character
Mean	: 0.03191	
3rd Qu.	: 5.02870	
Max.	: 11.17596	

Induce a binary decision tree (using `rpart`) and obtain the threshold value for the feature in the first split.

```

> data1$y=as.factor(data1$y)
> tree1=rpart(y~val,data1,method = "class")
> rpart.plot(tree1)
> printcp(tree1)

Classification tree:
rpart(formula = y ~ val, data = data1, method = "class")

variables actually used in tree construction:
[1] val


Root node error: 200/400 = 0.5

n= 400

   CP nsplit rel error xerror      xstd
1 0.99      0      1.00  1.140 0.0495076
2 0.01      1      0.01  0.015 0.0086277
> rpart.plot(tree1,main="Binary Tree For Dataset1")

```

How does this value compare to the empirical distribution of the feature? How many nodes does this tree have?



The Threshold value for features in the first split is 0.4. Tree has three nodes, one root, and two nodes leaf nodes. The above tree can classify both classes separately which shows empirical distribution.

What is the entropy and Gini at each?

```

> gini=function(p){
+ GI=2 * p* (1-p)
+ return (GI)
+ }
> entropy=function(p){
+ Entropy=(p * log(p)+(1-p)*log(1-p))
+ return(Entropy)
+ }
> |

```

```

> plot1=rpart.plot(tree1,extra = 106,fallen.leaves = TRUE,main="Binary Tree for Dataset1")
> plot2=rpart.plot(tree2,extra = 106,fallen.leaves = TRUE,main="Binary Tree for Dataset2")

```

```

> gini_index1=sapply(plot1$obj$frame$val2[,5],gini)
> gini_index1
[1] 0.50000000 0.00000000 0.01960592
> entropy1=sapply(plot1$obj$frame$val2[,5],entropy)
> entropy1
[1] -0.69314718      NaN -0.05554608
>

```

Repeat with normal distributions of (1,2) and (-1,2)

```

> set.seed(200)
> d3=data.frame(val=rnorm(n=200,mean=1,sd=2),y=rep("yes",200))
> d4=data.frame(val=rnorm(n=200,mean=-1,sd=2),y=rep("no",200))
> data2=rbind(d3,d4)
> summary(data2)
      val              y
Min.   :-6.43362    Length:400
1st Qu.:-1.38819    Class :character
Median : 0.09315    Mode  :character
Mean   : 0.03191
3rd Qu.: 1.44028
Max.   : 7.17596

```

Induce a binary decision tree (using rpart) and obtain the threshold value for the feature in the first split.

```

> data2$y=as.factor(data2$y)
> tree2=rpart(y~val,data2,method="class")
> printcp(tree2)

```

Classification tree:

```
rpart(formula = y ~ val, data = data2, method = "class")
```

Variables actually used in tree construction:

```
[1] val
```


Root node error: 200/400 = 0.5

n= 400

```

      CP nsplit rel error xerror      xstd
1 0.445    0    1.000   1.14 0.049508
2 0.020    1    0.555   0.64 0.046648
3 0.010    3    0.515   0.58 0.045376
> rpart.plot(tree2)

```



The threshold value for the first split is 0.62. The tree has total of 7 nodes, one root, and 4 leaf nodes. As normal distribution is the very near tree has many nodes with different labels in the node and has more overlapping of labels in nodes.

What are the entropy and Gini at each?

```

> gini_index2
[1] 0.5000000 0.4294896 0.2834393 0.4950000 0.4591837 0.4872000 0.3509353
> entropy2=sapply(plot2$obj$frame$yval2[,5],entropy)
> entropy2
[1] -0.6931472 -0.6208783 -0.4573738 -0.6881388 -0.6517566 -0.6802920 -0.5356184
>

```

Prune this tree (using part.prune) with a complexity parameter of 0.1. Describe the resulting tree.


```

> tree3=prune.rpart(tree2,cp=0.1)
> printcp(tree3)

Classification tree:
rpart(formula = y ~ val, data = data2, method = "class")

variables actually used in tree construction:
[1] val


Root node error: 200/400 = 0.5

n= 400

      CP nsplit rel error xerror      xstd
1 0.445      0   1.000   1.14 0.049508
2 0.100      1   0.555   0.64 0.046648
> plot3=rpart.plot(tree3,extra = 106,main="Purned Binary Tree")
> gini_index3=sapply(plot3$obj$frame$yval2[,5],gini)
> gini_index3
[1] 0.5000000 0.4294896 0.3509353
> entropy3=sapply(plot3$obj$frame$yval2[,5],entropy)
> entropy3
[1] -0.6931472 -0.6208783 -0.5356184
> |

```

Purned Binary Tree



The threshold value for the first split is 0.62. The tree has 3 nodes, one root node, and 2 leaf nodes. After the pruning process tree is much better than the previous one as this has only two leaf nodes with few overlapping labels.

Entropy and Gini value calculation:

Problem 2.2

Load the Wine Quality sample dataset from the UCI Machine Learning Repository (inequality-red.csv and winequality-white.csv) into R using a data frame.

```

> library(caret)
> library(randomForest)
> library(rpart)
> library(rpart.plot)
> white_winequality=read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv",header = TRUE,sep = ";",stringsAsFactors = TRUE)
> red_winequality=read.csv("https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv",header = TRUE,sep = ";",stringsAsFactors = TRUE)

```

```

> str(white_winequality)
'data.frame': 4898 obs. of 12 variables:
 $ fixed.acidity : num 7 6.3 8.1 7.2 7.2 8.3 8.2 7 6.3 8.1 ...
 $ volatile.acidity : num 0.27 0.3 0.28 0.23 0.23 0.28 0.32 0.27 0.3 0.22 ...
 $ citric.acid : num 0.36 0.34 0.4 0.32 0.32 0.4 0.16 0.36 0.34 0.43 ...
 $ residual.sugar : num 20.7 1.6 6.9 8.5 8.5 6.9 7 20.7 1.6 1.3 ...
 $ chlorides : num 0.045 0.049 0.05 0.058 0.058 0.05 0.045 0.045 0.049 0.044 ...
 $ free.sulfur.dioxide : num 45 14 30 47 47 30 30 45 14 28 ...
 $ total.sulfur.dioxide : num 170 132 97 180 186 97 138 170 132 139 ...
 $ density : num 1.001 0.998 0.995 0.996 0.998 ...
 $ pH : num 3.3 3.3 2.6 3.19 3.19 3.26 1.18 3.3 3.22 ...
 $ sulphates : num 0.45 0.49 0.44 0.4 0.4 0.44 0.47 0.45 0.49 0.45 ...
 $ alcohol : num 8.8 9.5 10.1 9.9 9.9 10.1 9.6 8.8 9.5 11 ...
 $ quality : int 8 6 6 6 6 6 6 6 6 ...

> summary(white_winequality)
fixed.acidity volatile.acidity citric.acid residual.sugar chlorides free.sulfur.dioxide total.sulfur.dioxide density pH
min.: 3.800 min.: 0.0800 min.: 0.0000 min.: 0.600 min.: 0.00900 min.: 2.00 min.: 9.0 min.: 0.8971 min.: 12.720
1st Qu.: 6.300 1st Qu.: 0.2100 1st Qu.: 0.2700 1st Qu.: 1.700 1st Qu.: 0.03600 1st Qu.: 23.00 1st Qu.: 108.0 1st Qu.: 0.9917 1st Qu.: 13.090
Median : 6.800 Median : 0.2000 Median : 0.3200 Median : 5.200 Median : 0.04100 Median : 34.00 Median : 134.0 Median : 0.9937 Median : 13.180
Mean : 6.851 Mean : 0.2782 Mean : 0.3542 Mean : 5.201 Mean : 0.04577 Mean : 35.31 Mean : 138.4 Mean : 0.9940 Mean : 13.188
3rd Qu.: 7.300 3rd Qu.: 0.3200 3rd Qu.: 0.3900 3rd Qu.: 9.900 3rd Qu.: 0.05000 3rd Qu.: 48.00 3rd Qu.: 167.0 3rd Qu.: 0.9981 3rd Qu.: 13.280
Max.: 14.200 Max.: 1.1000 Max.: 1.6600 Max.: 65.900 Max.: 0.34600 Max.: 289.00 Max.: 1440.0 Max.: 1.0390 Max.: 13.820

sulphates alcohol quality
min.: 0.2200 min.: 8.00 min.: 3.000
1st Qu.: 0.4100 1st Qu.: 9.30 1st Qu.: 5.000
Median : 0.4700 Median : 10.40 Median : 6.000
Mean : 0.4898 Mean : 10.31 Mean : 5.878
3rd Qu.: 0.5500 3rd Qu.: 11.40 3rd Qu.: 6.000
Max.: 1.0900 Max.: 14.20 Max.: 9.000

```

```

> str(red_winequality)
'data.frame': 1599 obs. of 12 variables:
 $ fixed.acidity : num 7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
 $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
 $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
 $ residual.sugar : num 1.9 2.6 2.3 1.9 1.9 1.8 1.8 1.2 2 6.1 ...
 $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
 $ free.sulfur.dioxide : num 11 23 15 17 11 13 13 15 9 17 ...
 $ total.sulfur.dioxide : num 34 67 54 60 34 40 39 21 18 102 ...
 $ density : num 0.998 0.997 0.997 0.998 0.998 ...
 $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
 $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.37 0.8 ...
 $ alcohol : num 9.4 9.8 9.8 9.6 9.4 9.4 9.4 10 9.5 10.3 ...
 $ quality : int 5 5 5 6 5 5 5 7 7 5 ...

> summary(red_winequality)
fixed.acidity volatile.acidity citric.acid residual.sugar chlorides free.sulfur.dioxide total.sulfur.dioxide density pH
min.: 4.60 min.: 0.1200 min.: 0.000 min.: 0.900 min.: 0.01200 min.: 1.00 min.: 6.00 min.: 0.9901 min.: 12.740
1st Qu.: 7.10 1st Qu.: 0.3900 1st Qu.: 0.090 1st Qu.: 1.900 1st Qu.: 0.07000 1st Qu.: 7.00 1st Qu.: 22.00 1st Qu.: 0.9956 1st Qu.: 13.210
Median : 7.90 Median : 0.5200 Median : 0.260 Median : 2.200 Median : 0.07900 Median : 14.00 Median : 38.00 Median : 0.9968 Median : 13.310
Mean : 8.32 Mean : 0.5278 Mean : 0.271 Mean : 2.559 Mean : 0.08747 Mean : 15.87 Mean : 48.47 Mean : 0.9967 Mean : 13.311
3rd Qu.: 9.20 3rd Qu.: 0.6400 3rd Qu.: 0.420 3rd Qu.: 2.600 3rd Qu.: 0.09000 3rd Qu.: 21.00 3rd Qu.: 62.00 3rd Qu.: 0.9978 3rd Qu.: 13.400
Max.: 15.90 Max.: 1.1500 Max.: 1.000 Max.: 15.500 Max.: 0.81100 Max.: 72.00 Max.: 1289.00 Max.: 1.0037 Max.: 14.010

sulphates alcohol quality
min.: 0.3100 min.: 8.40 min.: 3.000
1st Qu.: 0.5100 1st Qu.: 9.30 1st Qu.: 5.000
Median : 0.6200 Median : 10.20 Median : 6.000
Mean : 0.6581 Mean : 10.42 Mean : 5.836
3rd Qu.: 0.7300 3rd Qu.: 11.10 3rd Qu.: 6.000
Max.: 1.2000 Max.: 14.90 Max.: 9.000

```

Create an 80/20 test-train split of each wine data frame

```

> white_winequality$quality=as.factor(white_winequality$quality)
> index=createDataPartition(white_winequality$quality,p=0.8,list=FALSE)
> train_white=white_winequality[index,]
> test_white=white_winequality[-index,]
> dim(train_white)
[1] 3920 12
> dim(test_white)
[1] 978 12
> |

```

```

> red_winequality$quality=as.factor(red_winequality$quality)
> index1=createDataPartition(red_winequality$quality,p=0.8,list = FALSE)
> train_red=red_winequality[index1,]
> test_red=red_winequality[-index1,]
> dim(train_red)
[1] 1282 12
> dim(test_red)
[1] 317 12
>

```

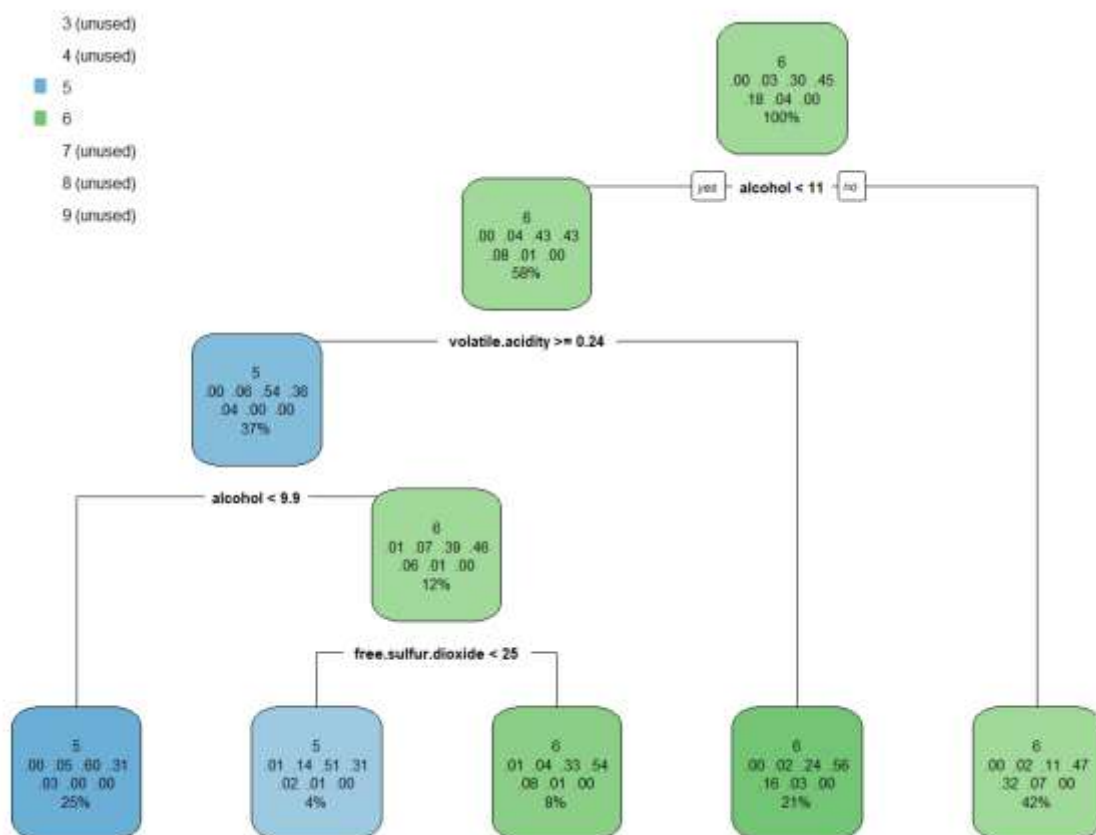
Decision Tree Model

Use the rpart package to induce a decision tree of both the red and white wines, targeting the quality output variable and visualize the tree using the part. Plot library

```

> tree_white=rpart(quality~.,data = train_white)
> rpart.plot(tree_white)

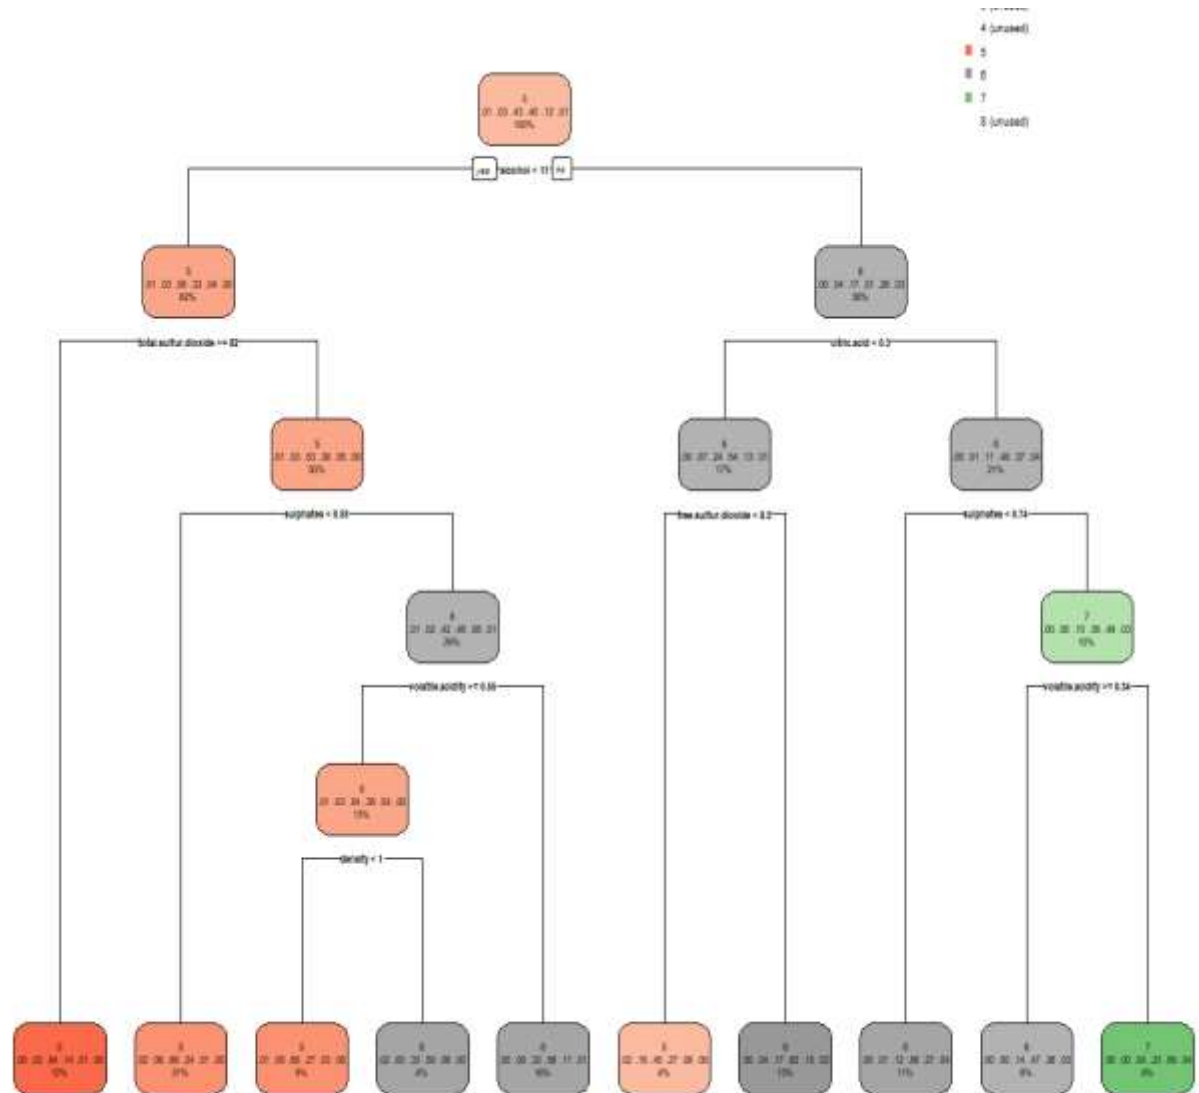
```



```

> tree_red=rpart(quantity~.,data = train_red)
> rpart.plot(tree_red)
>

```



use the caret package confusion Matrix method to determine the decision tree accuracy on the test set.

```
> tree_predict_white=predict(tree_white,test_white,type = "class")
> confusionMatrix(tree_predict_white,test_white$quality)
Confusion Matrix and Statistics
```

	Reference						
Prediction	3	4	5	6	7	8	9
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0
5	1	12	167	90	9	0	0
6	3	20	124	349	167	35	1
7	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0

overall statistics

Accuracy : 0.5276
 95% CI : (0.4958, 0.5593)
 No Information Rate : 0.4489
 P-value [Acc > NIR] : 4.729e-07

Kappa : 0.2051

McNemar's Test P-value : NA

Statistics by Class:

	Class: 3	Class: 4	Class: 5	Class: 6	Class: 7	Class: 8
Sensitivity	0.00000	0.00000	0.5739	0.7950	0.00	0.00000
Specificity	1.00000	1.00000	0.8370	0.3506	1.00	1.00000
Pos Pred Value	NaN	NaN	0.5986	0.4993	NaN	NaN
Neg Pred Value	0.99591	0.96728	0.8226	0.6774	0.82	0.96421
Prevalence	0.00409	0.03272	0.2975	0.4489	0.18	0.03579
Detection Rate	0.00000	0.00000	0.1708	0.3569	0.00	0.00000
Detection Prevalence	0.00000	0.00000	0.2853	0.7147	0.00	0.00000
Balanced Accuracy	0.50000	0.50000	0.7054	0.5728	0.50	0.50000

	Class: 9
Sensitivity	0.000000
Specificity	1.000000
Pos Pred Value	NaN
Neg Pred Value	0.998978
Prevalence	0.001022
Detection Rate	0.000000
Detection Prevalence	0.000000
Balanced Accuracy	0.500000

```
> |
```

```
> tree_predict_red=predict(tree_red,test_red,type = "class")
> confusionMatrix(tree_predict_red,test_red$quality)
Confusion Matrix and Statistics
```

	Reference						
Prediction	3	4	5	6	7	8	
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0
5	1	6	92	30	1	0	
6	1	4	43	89	31	2	
7	0	0	1	8	7	1	
8	0	0	0	0	0	0	0

```
Overall Statistics

          Accuracy : 0.5931
          95% CI   : (0.5367, 0.6476)
    No Information Rate : 0.429
    P-value [Acc > NIR] : 3.186e-09

          Kappa : 0.3247

  Mcnemar's Test P-value : NA

Statistics by Class:
```

	Class: 3	Class: 4	Class: 5	Class: 6	Class: 7	Class: 8
Sensitivity	0.000000	0.000000	0.6765	0.7008	0.17949	0.000000
Specificity	1.000000	1.000000	0.7901	0.5737	0.96403	1.000000
Pos Pred Value	NaN	NaN	0.7077	0.5235	0.41176	NaN
Neg Pred Value	0.993691	0.96845	0.7647	0.7415	0.89333	0.990536
Prevalence	0.006309	0.03155	0.4290	0.4006	0.12303	0.009464
Detection Rate	0.000000	0.000000	0.2902	0.2808	0.02208	0.000000
Detection Prevalence	0.000000	0.000000	0.4101	0.5363	0.05363	0.000000
Balanced Accuracy	0.500000	0.500000	0.7333	0.6372	0.57176	0.500000

```
> |
```

Compare the decision trees for red and white wine - what differences in terms of tree structure and variables of interest can be noted?

```
> table(tree_predict_white)
tree_predict_white
 3  4  5  6  7  8  9
0  0 279 699  0  0  0
> table(tree_predict_red)
tree_predict_red
 3  4  5  6  7  8
0  0 130 170 17  0
> |
```

```

> varImp(tree_white)
              overall
alcohol      205.048111
chlorides     88.026582
citric.acid   23.183313
density      114.509589
free.sulfur.dioxide 42.231833
pH           12.091773
residual.sugar  7.215579
total.sulfur.dioxide 57.891128
volatile.acidity 152.888927
fixed.acidity   0.000000
sulphates       0.000000
> varImp(tree_red)
              overall
alcohol      109.388331
chlorides     5.276332
citric.acid   21.855882
density       38.366207
fixed.acidity 34.513261
pH            4.443578
residual.sugar  3.356979
sulphates     76.659617
total.sulfur.dioxide 56.906256
volatile.acidity 87.244008
free.sulfur.dioxide 0.000000
> |

```

Use the Random Forest package to repeat the fit with a random forest tree model, and compare the resulting test accuracy against the original single tree model.

```

> rand_forest_white=randomForest(quality~.,data = train_white)
> rand_forest_predictw=predict(object = rand_forest_white,newdata = test_white)
> cm_rand_forest_white=confusionMatrix(data = rand_forest_predictw,reference = test_white$quality)
> rand_forest_red=randomForest(quality~.,data = train_red)
> rand_forest_predictr=predict(object = rand_forest_red,newdata = test_red)
> cm_rand_forest_red=confusionMatrix(data = rand_forest_predictr,reference = test_red$quality)

```

Original Tree Model Accuracy

```

> cm_red=confusionMatrix(tree_predict_red,test_red$quality)
> cm_red$overall["Accuracy"]
Accuracy
0.5930599
> cm_white=confusionMatrix(tree_predict_white,test_white$quality)
> cm_white$overall["Accuracy"]
Accuracy
0.5276074
> |

```

Rainforest Model Accuracy

```
> cm_rand_forest_red$overall["Accuracy"]
Accuracy
0.6940063
> cm_rand_forest_white$overall["Accuracy"]
Accuracy
0.6779141
> |
```

Random Forest performed better as it returned an accuracy of 69.4% for the red Wine Dataset. Random Forest returned an accuracy of 67.7% for the white Wine Dataset. The Accuracy increased from 59% to 69% in Random Forest Classifier in the red Wine Dataset and increased from 52% to 67% in Random Forest Classifier in the white Wine Dataset.

Problem 2.3

Load the SMS Spam Collection sample dataset from the UCI Machine Learning Repository (smsspamcollection.zip) into R using a data frame

```
> sms_data <- tempfile()
>
> download.file("https://archive.ics.uci.edu/ml/machine-learning-databases/00228/smsspamcollection.zip", sms_data)

> smsdata <- read.table(unz(sms_data, "SMSSpamCollection"), header = FALSE, sep = "\t", stringsAsFactors = FALSE)
```



```

> unlink(sms_data)
> colnames(smsdata) <- c("type","text")
>
> head(smsdata)
  type
1 ham
2 ham
3 spam
4 ham
5 ham
6 ham

                                text
1                                Go until jurong point,
  crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat...
2
                                ok lar... Joking wif u oni...
3                                Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 200
5. Text FA to 87121 to receive entry question(std txt rate)T&Cs apply 08452810075over18s
4
                                U dun say so early hor... U c already then say...
5 Nah I dont think he goes to usf, he lives around here though\nspam\iFreeMsg Hey there darling its been 3 weeks now and no wo
rd back! Id like some fun you up for it still? Tb ok! xxx std chgs to send, £1.50 to rcv
6
  Even my brother is not like to speak with me. They treat me like aids patent.
> str(smsdata)
'data.frame': 1779 obs. of 2 variables:
 $ type: chr "ham" "ham" "spam" "ham" ...
 $ text: chr "Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wa
t..." "ok lar... Joking wif u oni..." "Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to
receive entry question(") __truncated__ "U dun say so early hor... U c already then say..." ...
>

```

Use the tm package to create a Corpus of documents (Hint: Construct the corpus using a Vector Source of the text column)

```

> smsdata$type <- as.factor(smsdata$type)
> str(smsdata)
'data.frame': 1779 obs. of 2 variables:
 $ type: Factor w/ 2 levels "ham","spam": 1 1 2 1 1 1 1 2 2 1 ...
 $ text: chr "Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wa
t..." "ok lar... Joking wif u oni..." "Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to
receive entry question(") __truncated__ "U dun say so early hor... U c already then say..." ...
> summary(smsdata)
  type      text
ham :1541 Length:1779
spam: 238 Class :character
      Mode :character
>

```

```

> library(tm)
Loading required package: NLP
warning message:
package 'tm' was built under R version 4.2.2
> smsCorpus <- vCorpus(VectorSource(smsdata$text))
> print(smsCorpus)
<<vCorpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content:  documents: 1779
>

```

Inspect used to see a summary of a specific message

```
> inspect(head(smsCorpus,2))
<<VCorpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content: documents: 2

[[1]]
<<PlainTextDocument>>
Metadata: 7
Content: chars: 111

[[2]]
<<PlainTextDocument>>
Metadata: 7
Content: chars: 29

> |
```

as.character() used to view the desired message

```
> as.character(head(smsCorpus,1))
[1] "list(list(content = \"Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat...\", meta = list(author = character(0), timestamp = list(sec = 29.1702029705048, min = 36, hour = 5, mday = 1, mon = 10, year = 122, wday = 1, yday = 317, isdst = 0), description = character(0), heading = character(0), id = \"1\", language = \"en\", origin = character(0))))"
[2] "list()"

[3] "list()"
```

Apply the following transformations from the tm package to the corpus in order to prepare the data:

- a) Convert lowercase b) Remove stop words c) Strip whitespace d) Remove punctuation.

```
> smsdataT <- tm_map(smsCorpus, content_transformer(tolower))
>
> as.character(smsdataT[[1]])
[1] "go until jurong point, crazy.. available only in bugis n great world la e buffet... cine there got amore wat..."
> smsdataT <- tm_map(smsdataT, remove_words(stopwords()))
> as.character(smsdataT[[1]])
[1] "go jurong point, crazy.. available bugis n great world la e buffet... cine got amore wat..."
> smsdataT <- tm_map(smsdataT, strip_whitespace())
> as.character(smsdataT[[1]])
[1] "go jurong point, crazy.. available bugis n great world la e buffet... cine got amore wat..."
> smsdataT <- tm_map(smsdataT, remove_punctuation())
> as.character(smsdataT[[1]])
[1] "go jurong point crazy available bugis n great world la e buffet cine got amore wat"
> |
```

Use findFreqTerms to construct features from words occurring more than 10 times and proceed to split the data into a training and test set - for each creates a DocumentTermMatrix.

Create DocumentTermMatrix

Splitting into Train Test Split in 80% Training and 20% into Testing

```
> smsDTM <- DocumentTermMatrix(smsdataT)
> smsFreqterm <- findFreqTerms(smsDTM,lowfreq=10)
```

```
> head(smsFreqterm)
[1] "£100" "£1000" "£150" "£2000" "£250" "£350"
```

```
> smsfeatures.df <- as.data.frame(data.matrix(smsDTM),stringsAsFactors=FALSE)
> smsfeatures.df <- smsfeatures.df[,smsFreqterm]
> smsdataT.df <- cbind("Type" = smsdata$type, smsfeatures.df)
```

```
> library(caret)
> splitIndex <- createDataPartition(smsdataT.df$Type,p=0.8,list=FALSE)
>
> sms_train=data.matrix((smsdataT.df[splitIndex,-c(1)]))
> sms_test=data.matrix((smsdataT.df[-splitIndex,-c(1)]))
> sms_train_label=smsdataT.df[splitIndex,c(1)]
> sms_test_label=smsdataT.df[-splitIndex,c(1)]
> makeboolean=function(x){
+ x=ifelse(x>0,1,0)
+ }
> sms.train=apply(sms_train,2,makeboolean)
> sms.test=apply(sms_test,2,makeboolean)
> |
```

fit an SVM using the e1071 package.

```
> sms_clssifier=svm(sms.train,sms_train_label)
> summary(sms_clssifier)
```

```
Call:
svm.default(x = sms.train, y = sms_train_label)
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-kernel: radial
  cost: 1
```

```
Number of Support Vectors: 665
```

```
( 181 484 )
```

```
Number of Classes: 2
```

```
Levels:
ham spam
```

calculate test and train predictions, create confusion matrices

```
> sms_train_pred=predict(sms_clssifier,sms.train)
> sms_test_pred=predict(sms_clssifier,sms.test)
> sms_train_cm=confusionMatrix(sms_train_pred,sms_train_label)
> sms_test_cm=confusionMatrix(sms_test_pred,sms_test_label)
```

Training and Test Confusion Matrix

```
> sms_train_cm
Confusion Matrix and Statistics

          Reference
Prediction ham spam
ham      1233    21
spam       0   170

      Accuracy : 0.9853
      95% CI   : (0.9775, 0.9908)
No Information Rate : 0.8659
P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.9334

McNemar's Test P-Value : 1.275e-05

      Sensitivity : 1.0000
      Specificity : 0.8901
      Pos Pred Value : 0.9833
      Neg Pred Value : 1.0000
      Prevalence : 0.8659
      Detection Rate : 0.8659
      Detection Prevalence : 0.8806
      Balanced Accuracy : 0.9450

      'Positive' Class : ham

> |
```

```

> sms_test_cm
Confusion Matrix and Statistics

      Reference
Prediction ham spam
ham      308    21
spam      0     26

      Accuracy : 0.9408
      95% CI : (0.911, 0.963)
No Information Rate : 0.8676
P-value [Acc > NIR] : 5.803e-06

      Kappa : 0.6824

McNemar's Test P-value : 1.275e-05

      Sensitivity : 1.0000
      Specificity : 0.5532
      Pos Pred Value : 0.9362
      Neg Pred Value : 1.0000
      Prevalence : 0.8676
      Detection Rate : 0.8676
      Detection Prevalence : 0.9268
      Balanced Accuracy : 0.7766

      'Positive' class : ham
> |

```

Report your training and test set accuracy

```

> sms_train_cm$overall["Accuracy"]
Accuracy
0.9852528
> sms_test_cm$overall["Accuracy"]
Accuracy
0.9408451
> |

```